

Milestone 2: Using eBPF for Monitoring in Flexible Data Path

Team3

Harsh Khetawat, Nida Syed, Niklesh Phabiani, Sudhendu Kumar
{[hkhetaw](#), [nsyed](#), [nphabia](#), [skumar23](#)}@ncsu.edu

1. Motivation

There are tools like ftrace and perf_events that are part of the Linux kernel, which helps analyze Linux performance in detail. In that vein, eBPF was developed, first with low overhead latency histograms and extending to packet filtering capabilities. eBPF is now being integrated into the Linux kernel in stages. Its ability to be attached to kprobes is an attractive feature allows user defined programming on a live kernel image. And according to Brendan Gregg, eBPF will eventually be more than just a tracing tool.

2. Project objective (for Milestone 2)

- a. Understand the capabilities and advantages of eBPF based filtering
- b. Use eBPF to evaluate the measurement of our virtualized network topology along . We trace the packets going through our topology in different modes [Bridge (L2), Routed (L3), NAT, Tunnel] along with doing this in a containerized environment
- c. Packet header change for faster data path using TC

3. Related work and Literature Review

BPF or the Berkeley Packet Filter was designed for capturing network packets and filtering them based on specific rules^[1]. BPF is implemented as user-space programs that are run inside the kernel. However, the original BPF design did not take into account the advances of modern processors, such as the new instructions in multiprocessor systems.

The extended BPF (eBPF) was designed to overcome these limitations of BPF wherein the eBPF virtual machine resembles modern processors, which allows eBPF instructions to be mapped to the instruction set of the hardware for enhanced performance^[2]. This allows just-in-time(JIT) compilation. Additionally, an instruction makes it possible to call in-kernel functions without much cost. Also, the eBPF virtual machine is exposed directly to userspace. eBPF programs can be attached to tracepoints, kprobes, and perf events, which allows for tracing and analysing the network packets and the datapaths they take. There are several engineers who have used eBPF for packet tracing and analysis:

a. **Offloading OVS Flow Processing using eBPF:** William Tu from VMWare has used eBPF to offload OVS^[3]. The original OVS data path consists of receiving packets from the bridge/device hook with the OVS kernel datapath module performing the parsing, lookup with the server database, and required actions for each packet that arrives. Tu proposes replacing the entire OVS kernel datapath with eBPF, wherein ovs-vswitchd controls and manages the eBPF datapath. eBPF maps act as communication channels between the ovs-vswitchd and the eBPF datapath. In order to achieve this, Tu proceeds with parsing the header from P4, which is the OVS specification to an eBPF compiler specification from bcc. The flow table lookup involves the perf ring buffer which carries the packet and its metadata to the ovs-vswitchd which does flow translation, and program flow entry into the eBPF map. This allows actions such as flooding, mirror and push vlan, and tunnel, and the performance evaluation in terms of the number of packets sent is better as compared to the OVS datapath.

b. **XDP to implement load balancers, firewalls and other networking services (used in Cilium and other projects):** When it comes to high performance processing, XDP (eXpress Data Path) and TC (Traffic Control) are two hooks that are available. XDP provides a mechanism to run eBPF programs at the lowest level of the Linux networking stack, directly upon the receipt of the packet and immediately out of driver receive queues^[5]. XDP filters packets as soon as they are received using an eBPF program. This returns an action (XDP_PASS, XDP_DROP, XDP_TX, XDP_ABORTED), which can be used to define the context, etc.

While the kernel is in charge of moving packets quickly, BPF dictates the logic which decides the action. BPF provides the programming for XDP to allow users to access the low level networking data path to implement virtual switches, load balancers, firewalls, etc. XDP programs are also portable across different XDP platforms which includes not only Linux kernel but potentially NIC offloads, switches, DPDK, and other OSes.^[7]

Another technology initiated by Cisco that relies on BPF and XDP is Cilium. Cilium is an open source software that aims to provide network connectivity and load balancing functionality between workloads like containers^[13]. At the core of Cilium is eBPF which provides security and networking capabilities by the dynamic insertion of eBPF bytecode into the Linux kernel at various hooks or tracepoints. Some of the capabilities include core data path filtering, mangling, monitoring, and redirection. Cilium also uses XDP which enables special eBPF programs to run from the network driver with direct access to the packet's DMA buffer. This allows programs to be attached at the earliest possible point in the software stack allowing for a programmable, high-performance packet processing in the Linux kernel networking data path. Cilium also supports distributed load balancing for traffic between application containers and to external services. The load balancing is implemented using BPF through efficient hash tables allowing for scale, and supports direct server return if the load balancing operation is not performed on the source host.

Chen and Lu proposed a CETH(Common Ethernet Driver Framework) for XDP(eXpress Data Path) to overcome the kernel network performance for virtualization [9]. The

simplified CETH for XDP led to efficient memory and buffer management, customizable metadata structure, and compatibility with the kernel IP stack.

c. OVN:

Another great work related to eBPF is the IOVisor project. IOVisor is a community driven open source project that opened up new ways to innovate, develop and share I/O and networking functions. It allows creation of IOModules that can be used to build networking, security and tracing applications. It even provides a programmable data plane and development tools to simplify the creation and sharing of dynamic IOModules. Risso, Bertrone, and Bernal proposed combining IOVisor with the OVN(Open Virtual Network) architecture^[4]. OVN is an open source project to provide a standard for virtualized network switching functions^[7]. OVN provides L2 and L3 networking. Thus, OVN can be used to connect VMs or containers into private L2 and L3 networks quickly and programmatically without the use of VLANs or other physical network resources. OVN contains logical switches and routers, security groups, and L2,L3,L4 ACLs. It is implemented using a tunneled overlay network with protocols. OVN can be used with the OpenFlow protocol, which manipulates flow tables. OVN can also be used in OpenStack based networks, where Open vSwitch is the virtual switch. IOModule is an eBPF program that performs a specific task such as bridging, routing, NATing, etc. IO Modules can be combined to create complex services as in a service chain. The proposal was to keep OVN control plane and remove the OvS data plane, wherein this new IOVisor-OVN architecture reads information from the existing databases such as ovs-localDB(s), maps changes such as service requests into IOModules, and exploits Hover (REST API front-end) to inject, configure and bind IOModules to network interfaces^[4].

d. Creation of complex network services with eBPF:

Sebastiano Miano et al.^[8] exploited eBPF to create complex network functions, by considering the limitation of the size of the eBPF program, which is restricted to 4096 assembly instructions to guarantee that any program will terminate within a bounded amount of time. This restriction may be limiting when creating network functions that perform complex actions in the data plane. A learning for them because of this limitation was to partition the network function into multiple eBPF programs and jumping from one to another through tail calls. This technique enables the creations of network services as a collection of loosely coupled modules, each one implementing a different function. Also, network functions may have to put the current frame on hold while waiting for some other event to happen. An example is a router holding the packet while its ARP request is getting served. eBPF does not have the capability of holding the packet, thereby not allowing it to take ownership of the packet. A remedy suggested for this is the usage of a slow path module, executed in the userspace, that receives packets from the eBPF program and reacts consequently with respect to the processing defined by the developer. The necessity of such a slow path module is also highlighted in ^[9], wherein the authors use the OvS userspace module to process packets that do not match a flow in

the OvS kernel eBPF data path.

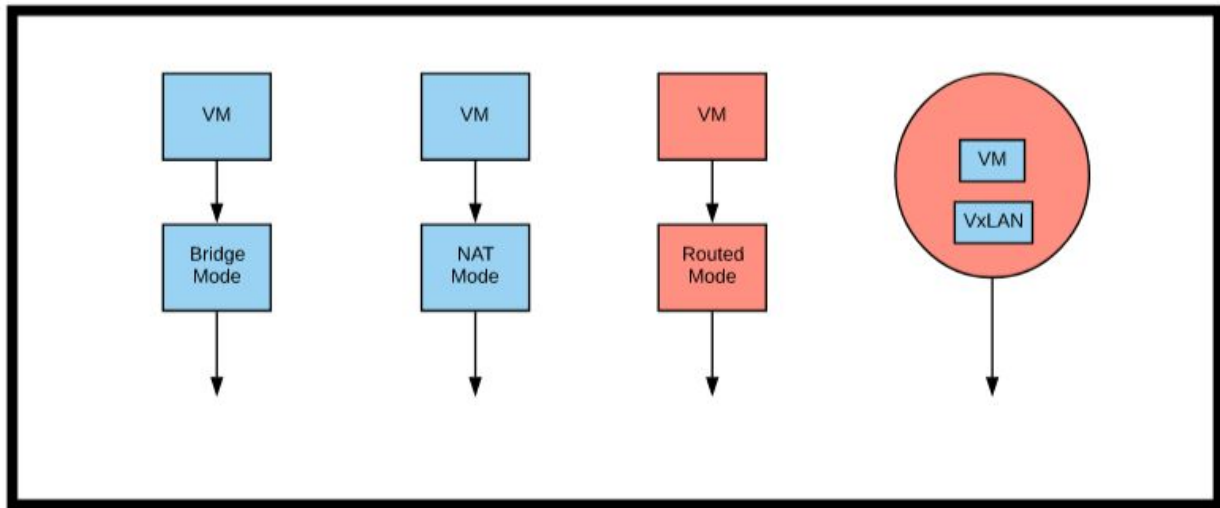
- e. **Flow sampling, monitoring:** Yunsong Lu implemented monitoring the whole virtual network stack from socket to virtual switch to physical NIC by making use of existing hooks which can be extended with eBPF and kprobes attached^[10]. By the use of Canal View (a container networking framework from Huawei), which is the topology-based virtual networking monitoring system, monitoring of Application-to-Application network traffic was done. This was implemented by making use of NFV(Network Functions Virtualization), Cloud Native, IOVisor, and E2. IOVisor was used as the MDPC (Micro Data Plane Container) with the applications developed and deployed with a Docker-like mechanism. This resulted in a faster network data plane in the kernel.
Now, when it comes to the usage of eBPF in Kubernetes, the Cilium project [16] explains how existing products (Cilium, Weave Scope) leverage eBPF to work with Kubernetes. In specific, it describes what interactions with eBPF are interesting in the context of container deployment.

4. Functional specs

- a. **Management Plane:** To be used as a **monitoring plugin** in a multi-tenant environment, the tenant provides the tenant ID to the plugin. The plugin-backend running on the hypervisor has information of tenant's network (bridges, network addresses). When the monitoring plugin starts, backend may receive information of other tenants as well at the tracepoints. The plugin-backend filters the only relevant tenant information and displays it on the front end.
- b. **Firewall / Fast Data Path:** **Firewall:** The plugin front-end accepts Tenant ID and Firewall rules in form of IPtables commands. Plugin backend evaluates if the rule can be applied on the TC (Traffic Control) of hypervisor without affecting other tenants. If yes, plugin backend puts and eBPF byte code on the hypervisor. Else, it simply puts the iptables rules as is. **Fast Data Path:** The plugin takes packet from tenant VM's interface. Instead of letting the packet go through *_xmit() function of all the devices, eBPF changes the header of packets and device_id which bypasses all the networking devices to go to the *_xmit() of hypervisor interface.

1. Results

a) Monitoring



Base Scenario: Evaluating the time spent by packet in each device

We have considered four topologies (Bridge mode, NAT mode, Routed mode and VxLAN tunnel mode). We have taken 10 batches of readings with each batch comprising of 100 iterations, and the average is calculated for each batch. This is done for each of the aforementioned topologies.

Input: Ping from VM to 8.8.8.8. For VxLAN device, Ping was done to the VM in another namespace (having the destination TEP).

Results:

The timestamp seen in the snapshots are in nanoseconds.

- **Case I (Bridge Mode):**

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/EBPF$ sudo python tracepkt.py
NETWORK NS      INTERFACE  TYPE  ADDRESSES      FUNCTION NAME      TIMESTAMP
[ 4026532201]    brovsns-ns request 192.168.0.10 -> 8.8.8.8 net_dev_queue    3371013232722767
[ 4026532201]    brovsns-ns request 192.168.0.10 -> 8.8.8.8 net_dev_start_xmit 3371013232739754
[ 4026531957]    brovsns-br request 192.168.0.10 -> 8.8.8.8 netif_rx    3371013232744839
[ 4026531957]    brovsns-br request 192.168.0.10 -> 8.8.8.8 net_dev_xmit 3371013232748073
[ 4026531957]    ovs-br request 192.168.0.10 -> 8.8.8.8 netif_rx    3371013233032755
[ 4026531957]    ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_queue    3371013233064420
[ 4026531957]    ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_start_xmit 3371013233071002
[ 4026531957]    ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_xmit 3371013233184599
[ 4026531957]    ens3 reply 8.8.8.8 -> 192.168.122.208 napi_gro_receive_entry 3371013241910688
[ 4026531957]    ovs-br reply 8.8.8.8 -> 192.168.0.10 net_dev_queue    3371013241938463
[ 4026531957]    ovs-br reply 8.8.8.8 -> 192.168.0.10 net_dev_start_xmit 3371013241942812
[ 4026531957]    ovs-br reply 8.8.8.8 -> 192.168.0.10 net_dev_xmit 3371013241959964
[ 4026531957]    brovsns-br reply 8.8.8.8 -> 192.168.0.10 net_dev_queue    3371013242065221
[ 4026531957]    brovsns-br reply 8.8.8.8 -> 192.168.0.10 net_dev_start_xmit 3371013242070855
[ 4026532201]    brovsns-ns reply 8.8.8.8 -> 192.168.0.10 netif_rx    3371013242074751
[ 4026532201]    brovsns-ns reply 8.8.8.8 -> 192.168.0.10 net_dev_xmit 3371013242077791
```


brovsns-ns and brovsns-br are the endpoints of the interface between VM and Bridge, and Ovs-br is the bridge.

Interface/ Device	1st	2nd	3rd	4th	5th	6th	7th	8th	9th	10th
Veth pair	7.731	7.915	7.881	7.688	8.511	8.118	7.635	8.275	8.470	7.892
Bridge	34.940	31.88	41.34	33.69	34.34	40.85	38.52	44.56	37.31	33.58
Veth pair	5.211	4.988	4.816	4.944	5.172	5.466	5.461	5.120	5.012	4.977
Hypervisor	234.80	250.63	234.34	230.45	238.65	224.28	222.53	218.80	253.41	232.695

All times are in microseconds

Bridge mode

Bridge mode



Maximum time is taken in the hypervisor. From our understanding, in addition to NATting, the hypervisor will even have to go through a much longer chain of iptables rules. In all the other cases too, maximum time is taken by the hypervisor.

```
[ 4026531957]      brovs-host  reply 8.8.8.8 -> 192.168.0.10 net_dev_queue 6898755419777435
[ 4026531957]      brovs-host  reply 8.8.8.8 -> 192.168.0.10 net_dev_start_xmit 6898755419785467
[ 4026531957]      brovs-br    reply 8.8.8.8 -> 192.168.0.10 netif_rx 6898755419791224
[ 4026531957]      brovs-br    reply 8.8.8.8 -> 192.168.0.10 net_dev_xmit 6898755419795923
[ 4026531957]      brovsns-br  reply 8.8.8.8 -> 192.168.0.10 net_dev_queue 6898755419847863
[ 4026531957]      brovsns-br  reply 8.8.8.8 -> 192.168.0.10 net_dev_start_xmit 6898755419853525
[ 4026532607]      brovsns-ns  reply 8.8.8.8 -> 192.168.0.10 netif_rx 6898755419858039
[ 4026532607]      brovsns-ns  reply 8.8.8.8 -> 192.168.0.10 net_dev_xmit 6898755419861876
[ 4026532607]      brovsns-ns  request 192.168.0.10 -> 8.8.8.8 net_dev_queue 6898756412242516
[ 4026532607]      brovsns-ns  request 192.168.0.10 -> 8.8.8.8 net_dev_start_xmit 6898756412274828
[ 4026531957]      brovsns-br  request 192.168.0.10 -> 8.8.8.8 netif_rx 6898756412283188
[ 4026531957]      brovsns-br  request 192.168.0.10 -> 8.8.8.8 net_dev_xmit 6898756412288398
[ 4026531957]      brovs-br    request 192.168.0.10 -> 8.8.8.8 net_dev_queue 6898756412308302
[ 4026531957]      brovs-br    request 192.168.0.10 -> 8.8.8.8 net_dev_start_xmit 6898756412313647
[ 4026531957]      brovs-host  request 192.168.0.10 -> 8.8.8.8 netif_rx 6898756412318677
[ 4026531957]      brovs-host  request 192.168.0.10 -> 8.8.8.8 net_dev_xmit 6898756412322697
[ 4026531957]      ens3      request 192.168.122.175 -> 8.8.8.8 net_dev_queue 6898756412360661
[ 4026531957]      ens3      request 192.168.122.175 -> 8.8.8.8 net_dev_start_xmit 6898756412371897
[ 4026531957]      ens3      request 192.168.122.175 -> 8.8.8.8 net_dev_xmit 6898756412564663
[ 4026531957]      ens3      reply 8.8.8.8 -> 192.168.122.175 napi_gro_receive_entry 6898756421163838
('veth pair: ', 8360L)
('Device: ', 30459L)
('veth pair: ', 5030L)
('Device: ', 245986L)
('Average with key ', 'brovsns-nsbrovsns-br', 7892L)
('Average with key ', 'brovsns-brbrovs-br', 33579L)
('Average with key ', 'brovs-brbrovs-host', 4977L)
('Average with key ', 'brovs-hostens3', 232695L)
```

Screenshot of the last iteration of a batch

- **Case II (NAT Mode):**


```

ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/EBPF$ sudo python tracepkt.py
[ 4026532433] natovsns-ns request 192.168.1.10 -> 8.8.8.8 net_dev_queue 3371369156944237
[ 4026532433] natovsns-ns request 192.168.1.10 -> 8.8.8.8 net_dev_start_xmit 3371369156957590
[ 4026531957] natovsns-nat request 192.168.1.10 -> 8.8.8.8 netif_rx 3371369156962586
[ 4026531957] natovsns-nat request 192.168.1.10 -> 8.8.8.8 net_dev_xmit 3371369156965838
[ 4026531957] ovs-nat request 192.168.1.10 -> 8.8.8.8 netif_rx 3371369157109016
[ 4026531957] ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_queue 3371369157145166
[ 4026531957] ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_start_xmit 3371369157151420
[ 4026531957] ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_xmit 3371369157256810
[ 4026531957] ens3 reply 8.8.8.8 -> 192.168.122.208 napi_gro_receive_entry 3371369165970247
[ 4026531957] ovs-nat reply 8.8.8.8 -> 192.168.1.10 net_dev_queue 3371369166002202
[ 4026531957] ovs-nat reply 8.8.8.8 -> 192.168.1.10 net_dev_start_xmit 3371369166007103
[ 4026531957] ovs-nat reply 8.8.8.8 -> 192.168.1.10 net_dev_xmit 3371369166023908
[ 4026531957] natovsns-nat reply 8.8.8.8 -> 192.168.1.10 net_dev_queue 3371369166124967
[ 4026531957] natovsns-nat reply 8.8.8.8 -> 192.168.1.10 net_dev_start_xmit 3371369166130417
[ 4026532433] natovsns-ns reply 8.8.8.8 -> 192.168.1.10 netif_rx 3371369166133785
[ 4026532433] natovsns-ns reply 8.8.8.8 -> 192.168.1.10 net_dev_xmit 3371369166136777

```


natovsns-ns and natovsns-br are the endpoints of the interface between VM and Bridge, and Ovs-br is the bridge.

Interface / Device	1st	2nd	3rd	4th	5th	6th	7th	8th	9th	10th
veth pair	8.224	8.273	8.162	8.750	9.041	7.932	8.392	8.516	8.731	8.046
NAT device	48.1	45.61	45.59	48.40	47.80	44.15	48.46	48.38	49.48	51.08
veth pair	4.884	4.679	4.714	4.889	5.143	4.531	4.850	5.012	4.801	5.042
Hypervisor	190.16	208.22	221.63	226.92	198.54	216.84	183.31	215.09	221.49	213.88

All times are in microseconds

NAT mode

NAT mode



- **Case III (Routed Mode):**

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/EBPF$ sudo python tracepkt.py
NETWORK NS      INTERFACE  TYPE  ADDRESSES      FUNCTION NAME      TIMESTAMP
[ 4026532493]    rovsns-ns1 request 192.168.2.10 -> 8.8.8.8 net_dev_queue 3375143740674908
[ 4026532493]    rovsns-ns1 request 192.168.2.10 -> 8.8.8.8 net_dev_start_xmit 3375143740693508
[ 4026532625]    rovsns-r1 request 192.168.2.10 -> 8.8.8.8 netif_rx 3375143740698862
[ 4026532625]    rovsns-r1 request 192.168.2.10 -> 8.8.8.8 net_dev_xmit 3375143740702057
[ 4026532625]    rovs-1 request 192.168.2.10 -> 8.8.8.8 net_dev_queue 3375143740718259
[ 4026532625]    rovs-1 request 192.168.2.10 -> 8.8.8.8 net_dev_start_xmit 3375143740721681
[ 4026531957]    rovs-2 request 192.168.2.10 -> 8.8.8.8 netif_rx 3375143740724493
[ 4026531957]    rovs-2 request 192.168.2.10 -> 8.8.8.8 net_dev_xmit 3375143740727030
[ 4026531957]    ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_queue 3375143740745718
[ 4026531957]    ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_start_xmit 3375143740751022
[ 4026531957]    ens3 request 192.168.122.208 -> 8.8.8.8 net_dev_xmit 3375143740854063
[ 4026531957]    ens3 reply 8.8.8.8 -> 192.168.122.208 napi_gro_receive_entry 3375143749578240
[ 4026531957]    rovs-2 reply 8.8.8.8 -> 192.168.2.10 net_dev_queue 3375143749606886
[ 4026531957]    rovs-2 reply 8.8.8.8 -> 192.168.2.10 net_dev_start_xmit 3375143749611469
[ 4026532625]    rovs-1 reply 8.8.8.8 -> 192.168.2.10 netif_rx 3375143749615123
[ 4026532625]    rovs-1 reply 8.8.8.8 -> 192.168.2.10 net_dev_xmit 3375143749618117
[ 4026532625]    rovsns-r1 reply 8.8.8.8 -> 192.168.2.10 net_dev_queue 3375143749645962
[ 4026532625]    rovsns-r1 reply 8.8.8.8 -> 192.168.2.10 net_dev_start_xmit 3375143749649992
[ 4026532493]    rovsns-ns1 reply 8.8.8.8 -> 192.168.2.10 netif_rx 3375143749652925
[ 4026532493]    rovsns-ns1 reply 8.8.8.8 -> 192.168.2.10 net_dev_xmit 3375143749655796
```


rovsns-ns and rovsns-br are the endpoints of the interface between VM and Bridge, and rovs-1 and rovs-2 is the interface of hanging on the hypervisor.

Interface / Device	1st	2nd	3rd	4th	5th	6th	7th	8th	9th	10th
Veth pair	9.241	8.853	7.872	8.734	8.328	8.645	8.061	8.515	8.217	8.085
Router NS	43.73	41.29	39.05	39.07	45.40	40.79	40.42	39.72	38.31	39.15
Veth pair	5.171	5.128	4.865	5.062	5.335	5.127	4.928	4.967	4.750	4.937
Hypervisor	258.93	221.63	211.06	213.11	221.21	235.97	187.22	213.92	238.77	226.35

All times are in microseconds

Routed mode

Routed mode



```
[ 4026532955] v-rovs-routers request 192.168.2.10 -> 8.8.8.8 net_dev_queue 6901279492330387
[ 4026532955] v-rovs-routers request 192.168.2.10 -> 8.8.8.8 net_dev_start_xmit 6901279492354280
[ 4026533108] v-routers-rovs request 192.168.2.10 -> 8.8.8.8 netif_rx 6901279492362720
[ 4026533108] v-routers-rovs request 192.168.2.10 -> 8.8.8.8 net_dev_xmit 6901279492368553
[ 4026533108] v-routers-host request 192.168.2.10 -> 8.8.8.8 net_dev_queue 6901279492395497
[ 4026533108] v-routers-host request 192.168.2.10 -> 8.8.8.8 net_dev_start_xmit 6901279492400907
[ 4026531957] v-host-routers request 192.168.2.10 -> 8.8.8.8 netif_rx 6901279492405979
[ 4026531957] v-host-routers request 192.168.2.10 -> 8.8.8.8 net_dev_xmit 6901279492410536
[ 4026531957] ens3 request 192.168.122.175 -> 8.8.8.8 net_dev_queue 6901279492439692
[ 4026531957] ens3 request 192.168.122.175 -> 8.8.8.8 net_dev_start_xmit 6901279492452559
[ 4026531957] ens3 request 192.168.122.175 -> 8.8.8.8 net_dev_xmit 6901279492592795
[ 4026531957] ens3 reply 8.8.8.8 -> 192.168.122.175 napi_gro_receive_entry 6901279501209612
('veth pair: ', 8440L)
('Device: ', 38187L)
('veth pair: ', 5072L)
('Device: ', 186816L)
('Average with key ', 'v-rovs-routersv-routers-rovs', 8085L)
('Average with key ', 'v-routers-rosv-routers-host', 39154L)
('Average with key ', 'v-routers-hostv-host-routers', 4937L)
('Average with key ', 'v-host-routersens3', 226353L)
[ 4026531957] v-host-routers reply 8.8.8.8 -> 192.168.2.10 net_dev_queue 6901279501268812
[ 4026531957] v-host-routers reply 8.8.8.8 -> 192.168.2.10 net_dev_start_xmit 6901279501277232
[ 4026533108] v-routers-host reply 8.8.8.8 -> 192.168.2.10 netif_rx 6901279501283592
[ 4026533108] v-routers-host reply 8.8.8.8 -> 192.168.2.10 net_dev_xmit 6901279501289381
[ 4026533108] v-routers-rovs reply 8.8.8.8 -> 192.168.2.10 net_dev_queue 6901279502652424
[ 4026533108] v-routers-rovs reply 8.8.8.8 -> 192.168.2.10 net_dev_start_xmit 6901279502667861
[ 4026532955] v-rovs-routers reply 8.8.8.8 -> 192.168.2.10 netif_rx 6901279502674314
[ 4026532955] v-rovs-routers reply 8.8.8.8 -> 192.168.2.10 net_dev_xmit 6901279502679297
```

Screenshot of the last iteration in a batch

- **Case IV (Tunnel Mode):**

	NETWORK NS	INTERFACE	TYPE	ADDRESSES	FUNCTION NAME	TIMESTAMP
[4026532699]	vxlan0	request	10.0.0.1 -> 10.0.0.2	net_dev_queue	3380054913635303
[4026532699]	vxlan0	request	10.0.0.1 -> 10.0.0.2	net_dev_start_xmit	3380054913643827
[4026532699]	veth0	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_queue	3380054913659638
[4026532699]	veth0	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_start_xmit	3380054913663134
[4026532553]	veth1	tunnel	192.168.100.1 -> 192.168.200.1	netif_rx	3380054913666962
[4026532553]	veth1	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_xmit	3380054913669916
[4026532553]	veth1	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_xmit	3380054913672750
[4026532553]	veth2	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_queue	3380054913686722
[4026532553]	veth2	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_start_xmit	3380054913689719
[4026532761]	veth3	tunnel	192.168.100.1 -> 192.168.200.1	netif_rx	3380054913692907
[4026532761]	veth3	tunnel	192.168.100.1 -> 192.168.200.1	net_dev_xmit	3380054913695431
[4026532761]	vxlan0	request	10.0.0.1 -> 10.0.0.2	napi_gro_receive_entry	3380054913894301
[4026532761]	vxlan0	reply	10.0.0.2 -> 10.0.0.1	net_dev_queue	3380054913919614
[4026532761]	vxlan0	reply	10.0.0.2 -> 10.0.0.1	net_dev_start_xmit	3380054913923458
[4026532761]	veth3	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_queue	3380054913931351
[4026532761]	veth3	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_start_xmit	3380054913934309
[4026532553]	veth2	tunnel	192.168.200.1 -> 192.168.100.1	netif_rx	3380054913937348
[4026532553]	veth2	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_xmit	3380054913940004
[4026532553]	veth2	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_xmit	3380054913942674
[4026532553]	veth1	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_queue	3380054913951971
[4026532553]	veth1	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_start_xmit	3380054913954724
[4026532699]	veth0	tunnel	192.168.200.1 -> 192.168.100.1	netif_rx	3380054913957361
[4026532699]	veth0	tunnel	192.168.200.1 -> 192.168.100.1	net_dev_xmit	3380054913959748
[4026532699]	vxlan0	reply	10.0.0.2 -> 10.0.0.1	napi_gro_receive_entry	3380054913969036

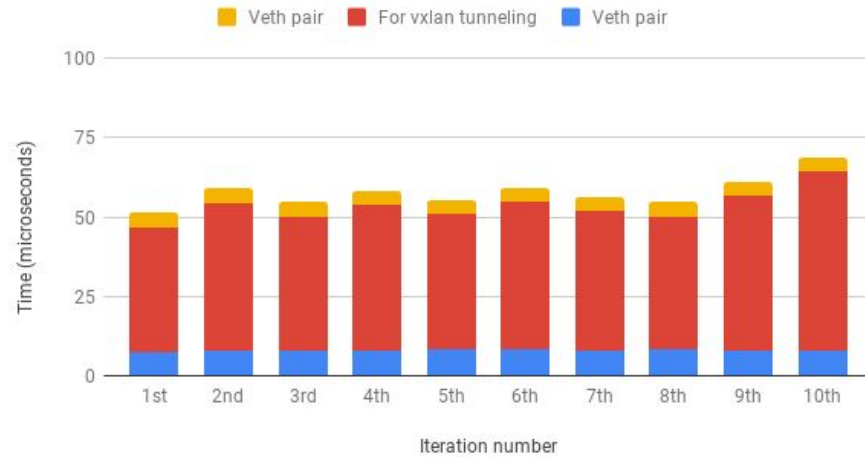
Vxlan0 is the src and destination VxLAN device. Veth0 - veth1 is the link between Src VM namespace and router (which is a namespace). veth2 - veth3 is the link between router and destination VM namespace. 10.0.0.1 and 10.0.0.2 are VMs separated by namespace router. 192.168.100.1 and 192.168.200.1 are the tunnel-endpoint IPs.

Interface / Device	1st	2nd	3rd	4th	5th	6th	7th	8th	9th	10th
Veth pair	7.416	8.038	7.815	7.811	8.147	8.239	7.947	8.126	7.973	7.986
For vxlan tunneling	39.47	46.23	42.35	45.94	42.73	46.54	43.99	41.98	48.74	56.29
Veth pair	4.475	5.016	4.564	4.423	4.618	4.559	4.248	4.70	4.577	4.582

All times are in microseconds

VxLAN tunnel mode


VxLAN tunnel mode



```
[ 4026533307] v-vxns1-br reply 192.168.4.1 -> 192.168.4.2 net_dev_xmit 6915827644253527
[ 4026533307] v-vxns1-br request 192.168.4.1 -> 192.168.4.2 net_dev_queue 6915827644253527
[ 4026533307] v-vxns1-br request 192.168.4.1 -> 192.168.4.2 net_dev_start_xmit 6915827644276544
[ 4026533590] v-br-vxns1 request 192.168.4.1 -> 192.168.4.2 netif_rx 6915827644285823
[ 4026533590] v-br-vxns1 request 192.168.4.1 -> 192.168.4.2 net_dev_xmit 6915827644292426
[ 4026533590] vxlan0 request 192.168.4.1 -> 192.168.4.2 net_dev_queue 6915827644322221
[ 4026533590] vxlan0 request 192.168.4.1 -> 192.168.4.2 net_dev_start_xmit 6915827644327571
[ 4026533664] vxlan0 request 192.168.4.1 -> 192.168.4.2 napi_gro_receive_entry 6915827644398700
[ 4026533664] v-br-vxns2 request 192.168.4.1 -> 192.168.4.2 net_dev_queue 6915827644413180
[ 4026533664] v-br-vxns2 request 192.168.4.1 -> 192.168.4.2 net_dev_start_xmit 6915827644417545
[ 402653376] v-vxns2-br request 192.168.4.1 -> 192.168.4.2 netif_rx 6915827644421854
[ 402653376] v-vxns2-br request 192.168.4.1 -> 192.168.4.2 net_dev_xmit 6915827644425974
[ 402653376] v-vxns2-br reply 192.168.4.2 -> 192.168.4.1 net_dev_queue 6915827644471860
('veth pair: ', 9279L)
('Device: ', 41748L)
('Device: ', 6915827644417545L)
('veth pair: ', 4309L)
('Average with key ', 'v-vxns1-brv-br-vxns1', 9513L)
('Average with key ', 'v-br-vxns1vxlan0', 57163L)
('Average with key ', 'vxlan0v-br-vxns2', 6915776971417655L)
('Average with key ', 'v-br-vxns2v-vxns2-br', 4957L)
[ 402653376] v-vxns2-br reply 192.168.4.2 -> 192.168.4.1 net_dev_start_xmit 6915827644477652
[ 4026533664] v-br-vxns2 reply 192.168.4.2 -> 192.168.4.1 netif_rx 6915827644482258
[ 4026533664] v-br-vxns2 reply 192.168.4.2 -> 192.168.4.1 net_dev_xmit 6915827644486007
[ 4026533664] vxlan0 reply 192.168.4.2 -> 192.168.4.1 net_dev_queue 6915827644496263
[ 4026533664] vxlan0 reply 192.168.4.2 -> 192.168.4.1 net_dev_start_xmit 6915827644501194
[ 4026533590] vxlan0 reply 192.168.4.2 -> 192.168.4.1 napi_gro_receive_entry 6915827644537151
[ 4026533590] v-br-vxns1 reply 192.168.4.2 -> 192.168.4.1 net_dev_queue 6915827644547489
[ 4026533590] v-br-vxns1 reply 192.168.4.2 -> 192.168.4.1 net_dev_start_xmit 6915827644552214
[ 4026533307] v-vxns1-br reply 192.168.4.2 -> 192.168.4.1 netif_rx 6915827644556217
[ 4026533307] v-vxns1-br reply 192.168.4.2 -> 192.168.4.1 net_dev_xmit 6915827644560025
```

Screenshot of the last iteration of a batch

Comparison between different modes



6. Future Work:

From our current topology we will use eBPF to evaluate the performance of network implementation of different Kubernetes CNI-plugins where we will measure the time spent by packet in OVS/linux bridges in different modes [Bridge (L2), Routed (L3), NAT, Tunnel]. We will also use eBPF with TC to Implement Fast Data Path/Firewall.

7. References:

- [*]<https://www.sdxcentral.com/sdn/network-virtualization/definitions/what-is-open-virtual-network-ovn-how-it-works/>
- [1] [A thorough introduction to eBPF](#)
- [2] [eBPF Based Networking](#)
- [3] [A presentation by William Tu of VMware on using eBPF to offload OVS: Offloading OVS flow processing using eBPF](#)
- [4] [A presentation by F. Risso, M. Bertrone, and M. Bernal on IO Visor and OVN: Coupling the flexibility of OVN with the efficiency of IOVisor: Architecture and Demo](#)
- [5] [A presentation by Jesper Brouer of Red Hat: XDP – Intro and Future Use Cases](#)
- [6] [High Speed Packet Filtering on Linux by Gilberto Bertin](#)
- [7] [Leveraging XDP for Programmable, High Performance Data Path in OpenStack](#)
- [8] Creation of complex network services with eBPF: [Creating Complex Network Services with eBPF: Experience and Lessons Learned](#)

- [9] [A presentation by Yunsong Lu and Yan Chen of Huawei on using IO Visor for Ethernet driver: CETH for XDP:Common Ethernet Driver Framework for Faster Network IO](#)
- [10] [A presentation by Yunsong Lu on using eBPF for virtual networking: Evolving Virtual Networking with IO Visor](#)
- [11] <https://github.com/cilium/cilium>
- [12] Building an Extensible Open vSwitch Datapath:
http://delivery.acm.org/10.1145/3140000/3139657/p72-tu.pdf?ip=152.7.255.201&id=3139657&acc=ACTIVE%20SERVICE&key=6ABC8B4C00F6EE47%2E4D4702B0C3E38B35%2E4D4702B0C3E38B35%2E4D4702B0C3E38B35&__acm__=1542139112_97364ee72f7825283ca39de5385c86b2
- [13] Why measure performance of CNI plugins:
<https://thenewstack.io/kubernetes-and-cni-whats-next-making-it-easier-to-write-networking-plugins/>
- [14] Feature comparison of various CNIs:
<https://chrislovecnm.com/kubernetes/cni/choosing-a-cni-provider/>
- [15] cgneta:
<https://github.com/kinvolk/cgneta>
- [16] Cilium:
<https://kubernetes.io/blog/2017/12/using-ebpf-in-kubernetes/>
- [17] Calico:
<https://www.projectcalico.org/wp-content/uploads/2018/04/ProjectCalico.v3.1.datasheet.pdf>
- [18] Flannel:
<https://github.com/coreos/flannel>
<https://coreos.com/flannel/docs/latest/flannel-config.html>
- [19] Weave:
<https://www.weave.works/docs/>
- [20] Understanding eBPF architecture:
https://sched.ws/hosted_files/kccnceu18/7f/KubeCon.pdf
- [21] eBPF based Container Networking:
<https://pdfs.semanticscholar.org/46d8/18dedcf1ed235a6f23ef826b02476897e5a1.pdf?ga=2.198904748.264721484.1538524519-1426312100.1538524519>
- [22] eBPF based Real-Time computing in IoT edge network:
<https://arxiv.org/pdf/1805.02797.pdf>
- [23] Tracing a packet's journey using Linux tracepoints, perf and eBPF:
<https://blog.yadutaf.fr/2017/07/28/tracing-a-packet-journey-using-linux-tracepoints-perf-e-bpf/>
<https://blog.yadutaf.fr/2016/03/30/turn-any-syscall-into-event-introducing-ebpf-kernel-probes/>
<https://jvns.ca/blog/2017/07/05/linux-tracing-systems/>

